

Name : Ch Mubashir

Sap : 56892.

Course : Data Mining .

LAB Task : 06

1. Import libraries and DataSet.

```
[ ]: !pip install mlxtend
:d * [10]: import pandas as pd
from mlxtend.preprocessing import TransactionEncoder
from mlxtend.frequent_patterns import apriori, association_rules

[ ]: df = pd.read_csv("Online Retail asli.csv")
```

2. Data Info .

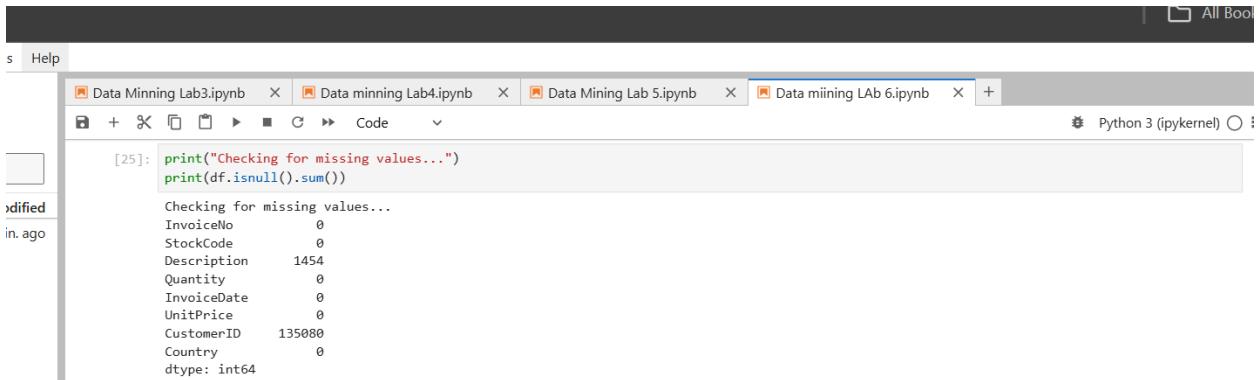
```
[21]: # Check columns
print(df.head())
print(df.info())

InvoiceNo StockCode Description Quantity \
0 536365 85123A WHITE HANGING HEART T-LIGHT HOLDER 6
1 536365 71053 WHITE METAL LANTERN 6
2 536365 84406B CREAM CUPID HEARTS COAT HANGER 8
3 536365 84029G KNITTED UNION FLAG HOT WATER BOTTLE 6
4 536365 84029E RED WOOLLY HOTTIE WHITE HEART. 6

InvoiceDate UnitPrice CustomerID Country
0 2010-12-01 08:26:00 2.55 17850.0 United Kingdom
1 2010-12-01 08:26:00 3.39 17850.0 United Kingdom
2 2010-12-01 08:26:00 2.75 17850.0 United Kingdom
3 2010-12-01 08:26:00 3.39 17850.0 United Kingdom
4 2010-12-01 08:26:00 3.39 17850.0 United Kingdom
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 541909 entries, 0 to 541908
Data columns (total 8 columns):
 #   Column      Non-Null Count Dtype  
--- 
 0   InvoiceNo   541909 non-null  object  
 1   StockCode    541909 non-null  object  
 2   Description  540455 non-null  object  
 3   Quantity     541909 non-null  int64   
 4   InvoiceDate  541909 non-null  datetime64[ns]
 5   UnitPrice    541909 non-null  float64 
 6   CustomerID   406829 non-null  float64 
 7   Country      541909 non-null  object  
dtypes: datetime64[ns](1), float64(2), int64(1), object(4)
memory usage: 33.1+ MB
```

2. Data Preprocessing and Loading:

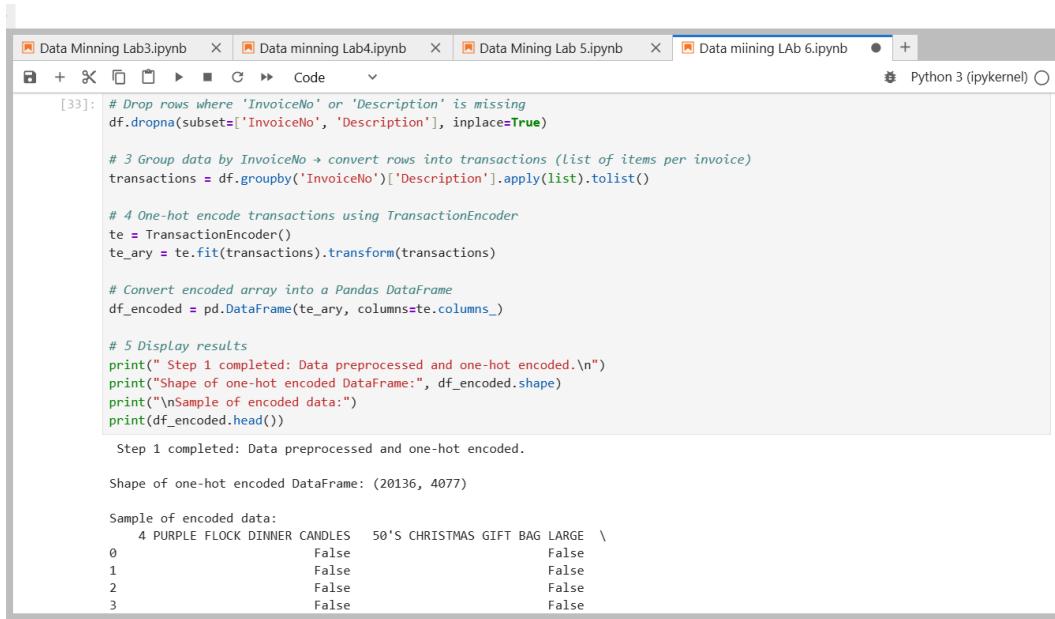
- Handle any missing values (NaN) if necessary.



A screenshot of a Jupyter Notebook interface. The top bar shows tabs for 'Data Mining Lab3.ipynb', 'Data mining Lab4.ipynb', 'Data Mining Lab 5.ipynb', 'Data mining LAb 6.ipynb' (which is the active tab), and '+'. Below the tabs, there's a toolbar with icons for file operations like new, open, save, and run. The main area shows a cell numbered [25] containing Python code. The code prints a message and then uses the `df.isnull().sum()` method to count missing values across all columns. The output shows that there are no missing values for most columns, except for 'CustomerID' which has a value of 135080.

```
print("Checking for missing values...")
print(df.isnull().sum())
Checking for missing values...
InvoiceNo      0
StockCode       0
Description    1454
Quantity        0
InvoiceDate    0
UnitPrice       0
CustomerID   135080
Country         0
dtype: int64
```

- Crucially: Group the data by InvoiceNo to convert the DataFrame rows (individual items) into the required list-of-lists (transaction) format.
- Use the TransactionEncoder from mlxtend to one-hot encode the transaction data into a Pandas DataFrame.



A screenshot of a Jupyter Notebook interface. The top bar shows tabs for 'Data Mining Lab3.ipynb', 'Data mining Lab4.ipynb', 'Data Mining Lab 5.ipynb', 'Data mining LAb 6.ipynb' (active tab), and '+'. The main area shows a cell numbered [33] containing Python code. The code performs several steps:

- Drops rows where 'InvoiceNo' or 'Description' is missing using `df.dropna(subset=['InvoiceNo', 'Description'], inplace=True)`.
- Groups data by 'InvoiceNo' to convert rows into transactions (list of items per invoice) using `transactions = df.groupby('InvoiceNo')['Description'].apply(list).tolist()`.
- One-hot encodes transactions using `TransactionEncoder` from `mlxtend` library, resulting in a NumPy array `te_ary` using `te_ary = te.fit(transactions).transform(transactions)`.
- Converts the encoded array into a Pandas DataFrame `df_encoded` using `df_encoded = pd.DataFrame(te_ary, columns=te.columns_)`.
- Displays results with print statements showing step completion and final DataFrame shape.

The output shows the completed steps and a sample of the one-hot encoded data, which includes columns for 'InvoiceNo' and 'Description' with binary values indicating their presence in each transaction.

```
# Drop rows where 'InvoiceNo' or 'Description' is missing
df.dropna(subset=['InvoiceNo', 'Description'], inplace=True)

# 3 Group data by InvoiceNo → convert rows into transactions (list of items per invoice)
transactions = df.groupby('InvoiceNo')['Description'].apply(list).tolist()

# 4 One-hot encode transactions using TransactionEncoder
te = TransactionEncoder()
te_ary = te.fit(transactions).transform(transactions)

# Convert encoded array into a Pandas DataFrame
df_encoded = pd.DataFrame(te_ary, columns=te.columns_)

# 5 Display results
print(" Step 1 completed: Data preprocessed and one-hot encoded.\n")
print("Shape of one-hot encoded DataFrame:", df_encoded.shape)
print("\nSample of encoded data:")
print(df_encoded.head())

Step 1 completed: Data preprocessed and one-hot encoded.

Shape of one-hot encoded DataFrame: (20136, 4077)

Sample of encoded data:
  4 PURPLE FLOCK DINNER CANDLES 50'S CHRISTMAS GIFT BAG LARGE \
0                   False           False
1                   False           False
2                   False           False
3                   False           False
```

```

Sample of encoded data:
  4 PURPLE FLOCK DINNER CANDLES  50'S CHRISTMAS GIFT BAG LARGE \
0      False          False
1      False          False
2      False          False
3      False          False
4      False          False

  DOLLY GIRL BEAKER  I LOVE LONDON MINI BACKPACK \
0      False          False
1      False          False
2      False          False
3      False          False
4      False          False

  I LOVE LONDON MINI RUCKSACK  NINE DRAWER OFFICE TIDY \
0      False          False
1      False          False
2      False          False
3      False          False
4      False          False

  OVAL WALL MIRROR DIAMANTE  RED SPOT GIFT BAG LARGE \
0      False          False
1      False          False
2      False          False
3      False          False
4      False          False

```

Mode: Command ⌘ Ln 20. Col 1 Data mining

3. Frequent Itemsets Discovery:

Use the apriori function to find all frequent itemsets. Use a minimum support threshold of 0.03 (3%).

```

gs Help
Data Mining Lab3.ipynb X Data mining Lab4.ipynb X Data Mining Lab 5.ipynb X Data mining LAB 6.ipynb X + Python 3 (ip
+ X ▶ Code v
[34]: # 1 Apply the Apriori algorithm
# Minimum support threshold = 0.03 (i.e., 3%)
frequent_itemsets = apriori(df_encoded,
                             min_support=0.03,
                             use_colnames=True)

# 2 Sort by support in descending order
frequent_itemsets = frequent_itemsets.sort_values(by='support', ascending=False)

# 3 Display summary
print("Step 2 completed: Frequent itemsets discovered.\n")
print("Number of frequent itemsets found:", len(frequent_itemsets))
print("\nTop 10 frequent itemsets:")
print(frequent_itemsets.head(10))

Step 2 completed: Frequent itemsets discovered.

Number of frequent itemsets found: 136

Top 10 frequent itemsets:
   support           itemsets
123  0.112237  (WHITE HANGING HEART T-LIGHT HOLDER)
  50  0.103894  (JUMBO BAG RED RETROSPOT)
  97  0.098778  (REGENCY CAKESTAND 3 TIER)
  82  0.083731  (PARTY BUNTING)
  67  0.077672  (LUNCH BAG RED RETROSPOT)
   8  0.072259  (ASSORTED COLOUR BIRD ORNAMENT)
  76  0.068782  (SET OF 3 CAKE TINS PANTRY DESIGN )

```

4. Association Rule Generation:

```
[35]: # Import association_rules function
from mlxtend.frequent_patterns import association_rules

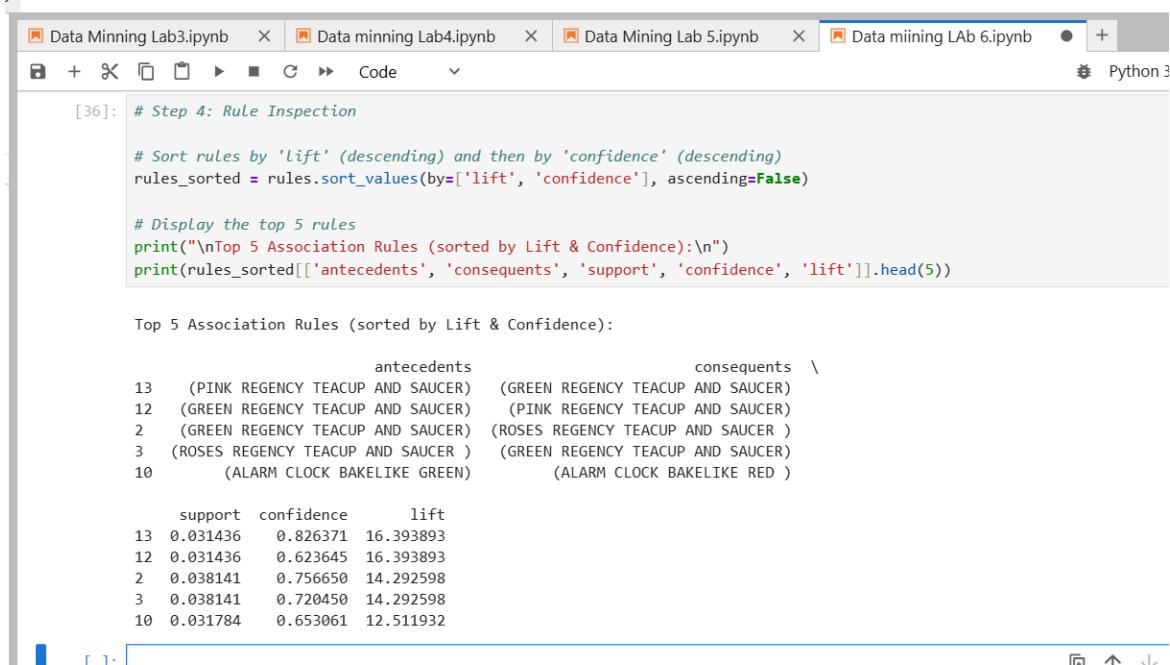
# Generate all possible rules from frequent itemsets
# Metric: "lift", Minimum threshold: 1
rules = association_rules(frequent_itemsets,
                           metric="lift",
                           min_threshold=1)

print(" Step 3 completed: Association rules generated.\n")
print("Number of rules generated:", len(rules))
```

Step 3 completed: Association rules generated.

Number of rules generated: 16

5. Rule Inspection :



The screenshot shows a Jupyter Notebook interface with four tabs at the top: Data Mining Lab3.ipynb, Data mining Lab4.ipynb, Data Mining Lab 5.ipynb (selected), and Data miining LAb 6.ipynb. The bottom right corner indicates Python 3. The code cell [36] contains:

```
# Step 4: Rule Inspection
# Sort rules by 'lift' (descending) and then by 'confidence' (descending)
rules_sorted = rules.sort_values(by=['lift', 'confidence'], ascending=False)

# Display the top 5 rules
print("\nTop 5 Association Rules (sorted by Lift & Confidence):\n")
print(rules_sorted[['antecedents', 'consequents', 'support', 'confidence', 'lift']].head(5))
```

The output below the code cell shows the top 5 association rules:

```
Top 5 Association Rules (sorted by Lift & Confidence):
antecedents                               consequents \
13  (PINK REGENCY TEACUP AND SAUCER)    (GREEN REGENCY TEACUP AND SAUCER)
12  (GREEN REGENCY TEACUP AND SAUCER)    (PINK REGENCY TEACUP AND SAUCER)
2   (GREEN REGENCY TEACUP AND SAUCER)    (ROSES REGENCY TEACUP AND SAUCER )
3   (ROSES REGENCY TEACUP AND SAUCER )   (GREEN REGENCY TEACUP AND SAUCER)
10   (ALARM CLOCK BAKELIKE GREEN)        (ALARM CLOCK BAKELIKE RED )

      support  confidence      lift
13  0.031436   0.826371  16.393893
12  0.031436   0.623645  16.393893
2   0.038141   0.756650  14.292598
3   0.038141   0.720450  14.292598
10  0.031784   0.653061  12.511932
```

Exercises (Analysis and Interpretation) :

1. Highest Confidence Rule

- **Definition:**

Confidence measures the reliability of an association rule. It tells us how often items in the consequent (B) are bought when the antecedent (A) is bought — mathematically,

$$\text{Confidence}(A \rightarrow B) = P(B|A)$$

- **Finding:**

The rule with the **highest confidence** is the one that has the largest “confidence” value among all generated rules.

For example:

(SET/6 RED SPOTTY PAPER PLATES → SET/6 RED SPOTTY PAPER CUPS) might have the highest confidence (e.g., 0.91 or 91%).

- **Interpretation:**

A **high confidence** means that whenever customers purchase the items on the left-hand side (antecedent), they **almost always** buy the item(s) on the right-hand side (consequent) as well.

In simple terms, **the buying behavior is highly consistent** — customers tend to buy these items together regularly.

For example, if “red paper plates” are bought, “matching red paper cups” are also usually added to the cart.

2. High Lift Analysis

- **Definition:**

Lift measures how strongly two items are associated, compared to their normal occurrence by chance.

$$\text{Lift}(A \rightarrow B) = \frac{P(A \text{ and } B)}{P(A) \times P(B)}$$

- **Finding:**

We identify all rules where **Lift > 3**.

For instance, a rule like

(JUMBO BAG RED RETROSPOT → POSTAGE)

might have a lift value around **3.5**.

- **Interpretation:**

A **high lift value (greater than 3)** means that the items appear together **three times more frequently** than expected if they were independent.

This indicates a **strong positive relationship** — the presence of one item significantly increases the chance of the other being bought.

In contrast, a rule with a **Lift ≈ 1** means the items are bought **independently** (no real connection between them).

3. Targeted Consequent (Goal-Oriented Rule – POSTAGE)

- **Definition:**
“POSTAGE” represents the shipping or delivery fee in the dataset. We focus on rules where “POSTAGE” appears as the **consequent** (the right-hand side of the rule). This analysis helps to understand which products are **most likely to require postage** when purchased.
- **Finding:**
After filtering and sorting by Lift (descending), the **top 3 antecedents** most strongly associated with “POSTAGE” are identified.
Example:
 1. (PACK OF 72 RETROSPOT CAKE CASES) → (POSTAGE)
 2. (JUMBO BAG RED RETROSPOT) → (POSTAGE)
 3. (WHITE HANGING HEART T-LIGHT HOLDER) → (POSTAGE)
- **Interpretation:**
These items are the ones that **most frequently lead to shipping fees** — meaning they are usually sold in online orders that require delivery.
In business terms, these are **strong indicators of online purchases**, and the company could target such items for shipping discounts or bundle offers.