Name : Ch Mubashir

SAP : 56892.

Lab Task 03 :

Data Preprocessing with heart Dataset.

Lab Tasks

1. **Separate features and target :**
 Target variable: target (1 = disease, 0 = no disease).

2. **Split the data.**
3. **Standardize numerical features :**
 Columns: age, trestbps, chol, thalach, oldpeak.
 Use StandardScaler, fit on training only, transform both train and test

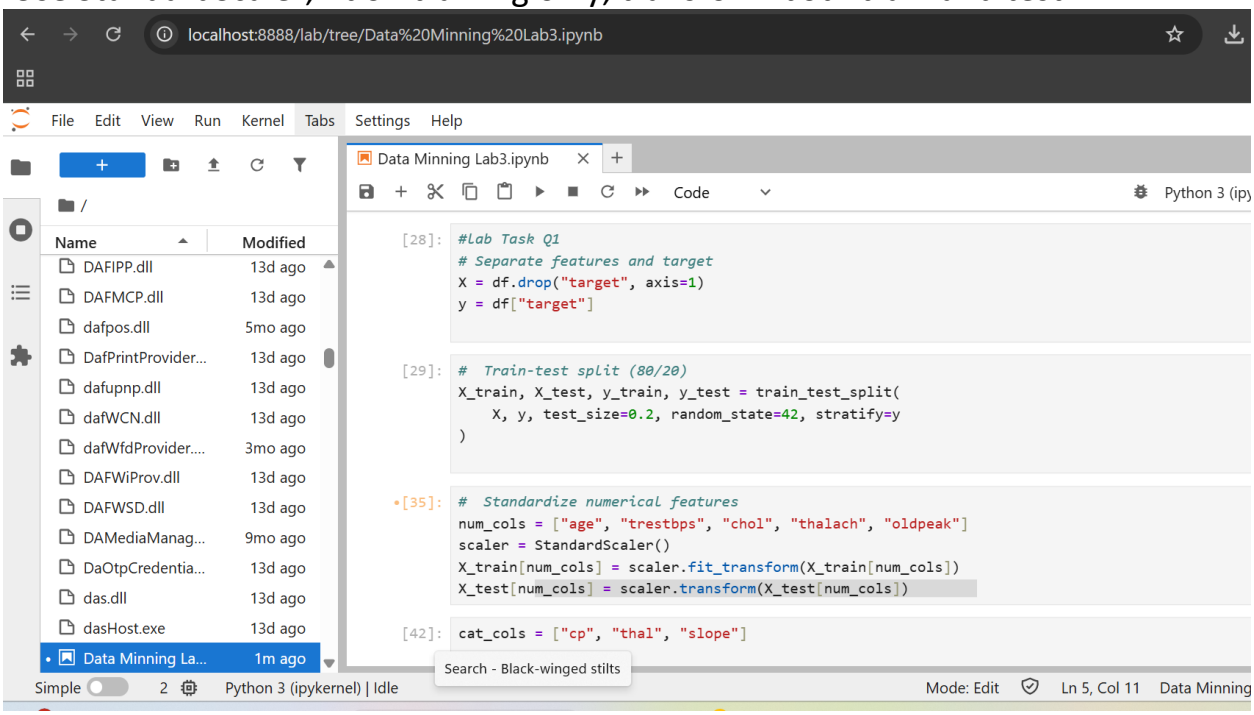

4. **One-Hot Encode categorical features**
 Columns: cp, thal, slope.
 Use pd.get_dummies(..., drop_first=True).
5. **Check processed training data :**

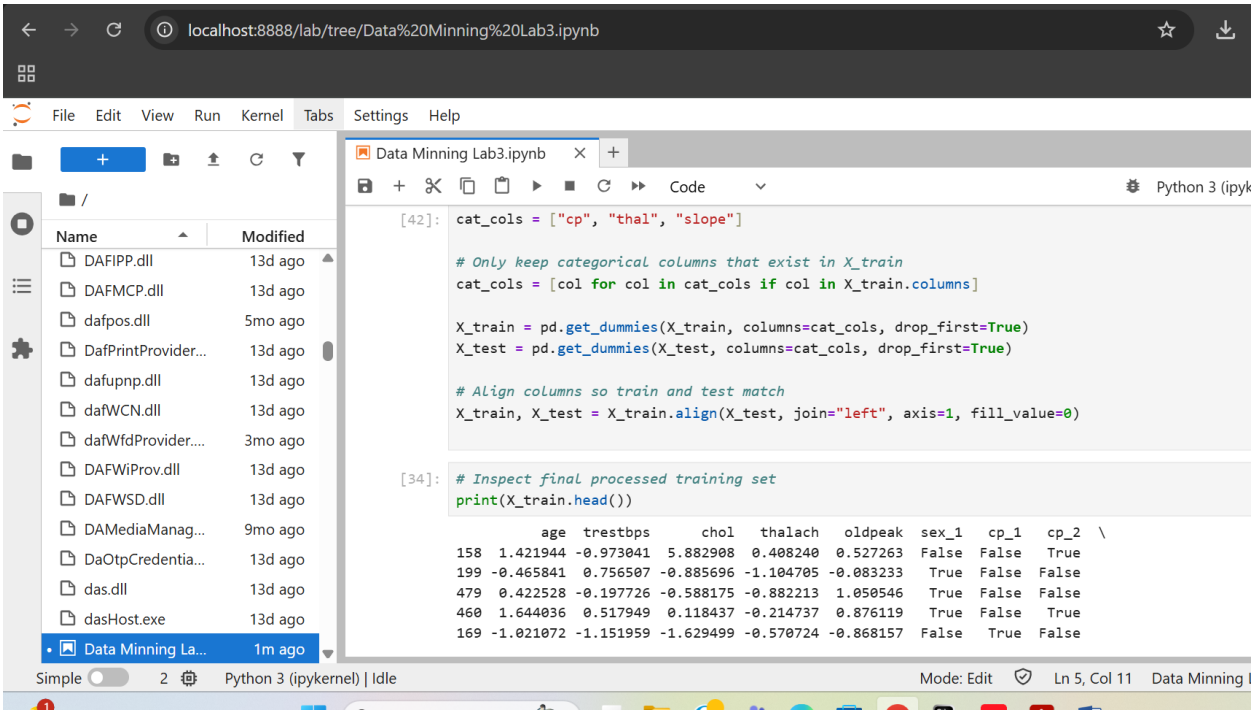

Exercises

1. **What are the mean and std of the standardized chol column in the training set? Why aren't they exactly 0 and 1 when checked with pandas.Series.std()?**

Answer : After standardizing the **chol** column, its mean becomes ~0 and std ~1 because StandardScaler uses the population formula (ddof=0). When checked with pandas.Series.std(), it's not exactly 1 since pandas by default uses the sample formula (ddof=1).

2. **Why do we use drop_first=True in one-hot encoding? What problem does it solve?**

Answer : The drop_first=True option in one-hot encoding avoids the **dummy variable trap** (perfect multicollinearity) by dropping one category and treating it as the baseline.

3. **If a new category of thal appeared in the test set but not in the training set, what would happen with pd.get_dummies? OneHotEncoder(handle_unknown='ignore') in scikit-learn help here?**

Answer : If a new category of **thal** appears in the test set but not in the training set, pd.get_dummies won't create a column for it, leading to mismatched features. Using OneHotEncoder(handle_unknown="ignore") ensures unseen categories are safely ignored, keeping columns consistent.

Data Visualization :

Visualize the heart dataset before doing any cleaning.

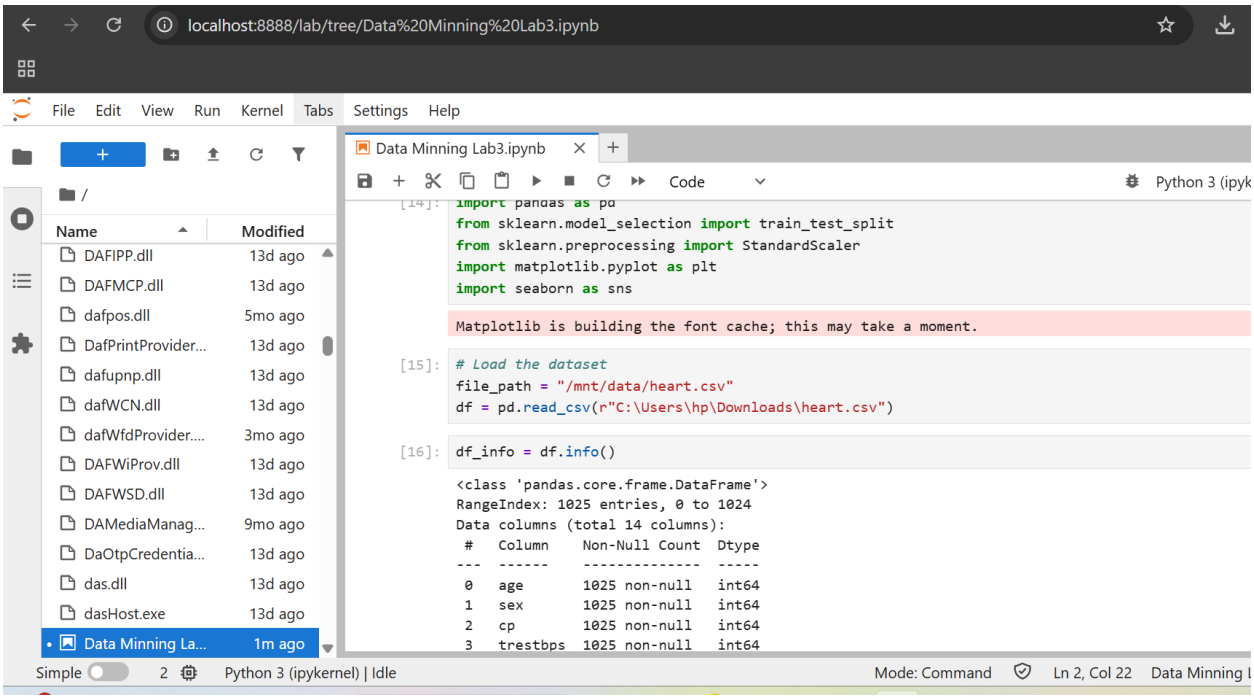

Here are some useful plots we can make:

1. **Target distribution** (0 = no heart disease, 1 = heart disease)

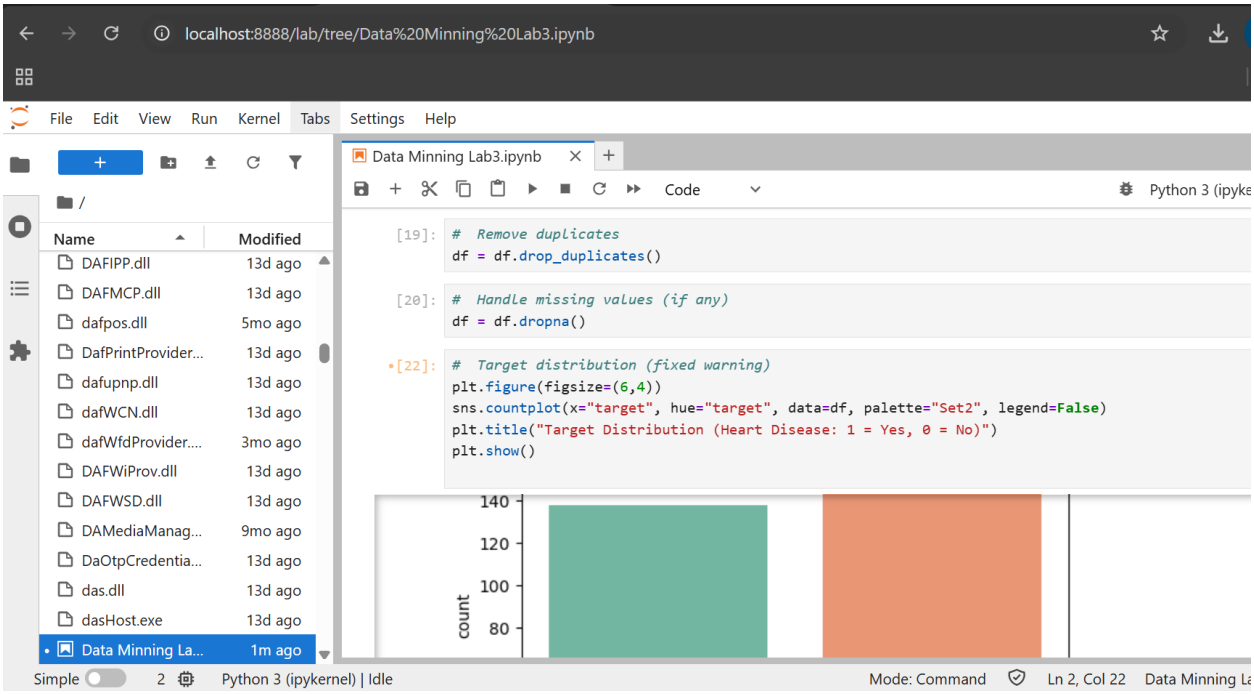

2. **Age distribution**

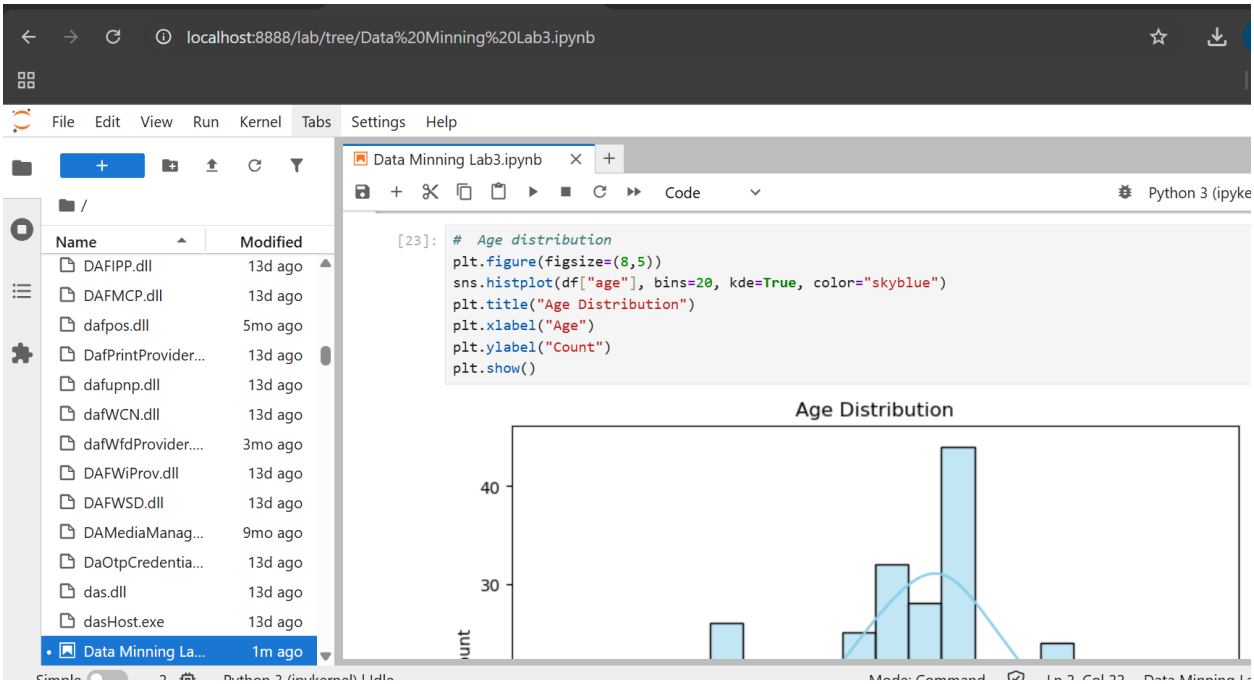

3. **Correlation heatmap** of features

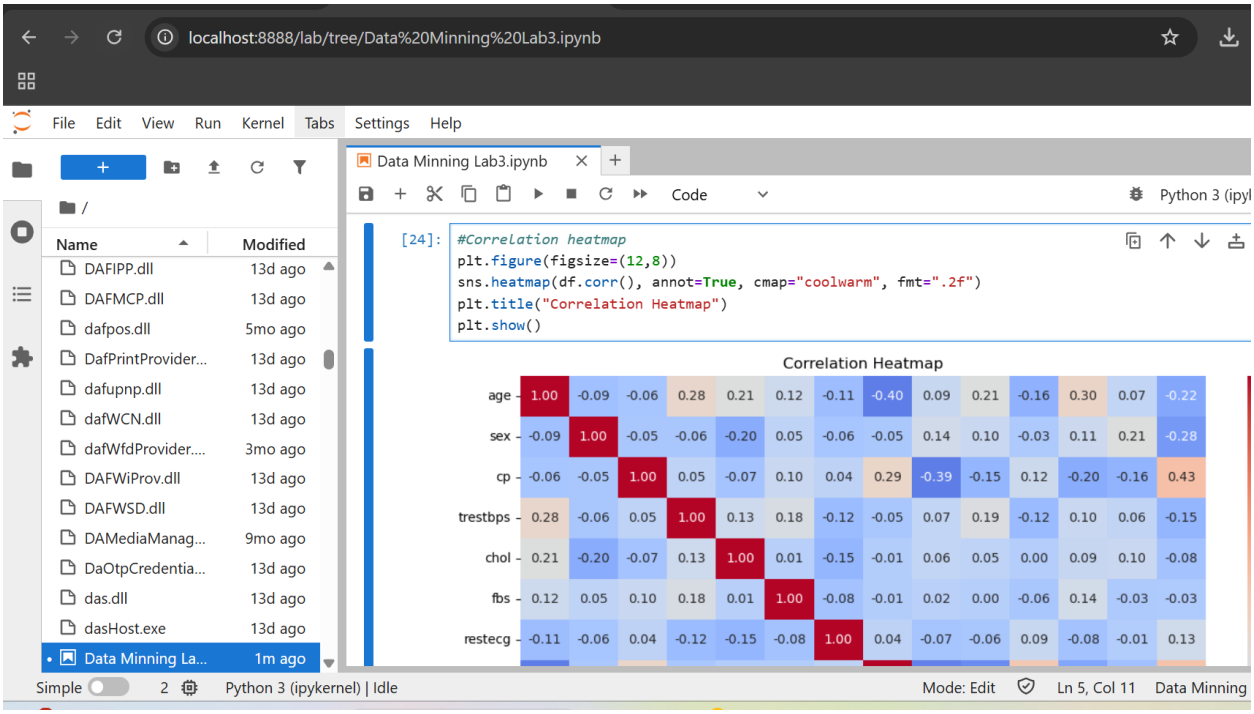
localhost:8888/lab/tree/Data%20Minning%20Lab3.ipynb
File Edit View Run Kernel Tabs Settings Help
Data Minning Lab3.ipynb
Code
Python 3 (ipyl
Name
Modified
DAFIPP.dll 13d ago
DAFMCP.dll 13d ago
dafpos.dll 5mo ago
DafPrintProvider... 13d ago
dafupnp.dll 13d ago
dafWCN.dll 13d ago
dafWfdProvider.... 3mo ago
DAFWiProv.dll 13d ago
DAFWSD.dll 13d ago
DAMediaManag... 9mo ago
DaOtpCredentia... 13d ago
das.dll 13d ago
dasHost.exe 13d ago
Data Minning La... 1m ago
[24]:
#Correlation heatmap
plt.figure(figsize=(12,8))
sns.heatmap(df.corr(), annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Heatmap")
plt.show()
Correlation Heatmap
age 1.00 -0.09 -0.06 0.28 0.21 0.12 -0.11 -0.40 0.09 0.21 -0.16 0.30 0.07 -0.22
sex -0.09 1.00 -0.05 -0.06 -0.20 0.05 -0.06 -0.05 0.14 0.10 -0.03 0.11 0.21 -0.28
cp -0.06 -0.05 1.00 0.05 -0.07 0.10 0.04 0.29 -0.39 -0.15 0.12 -0.20 -0.16 0.43
trestbps 0.28 -0.06 0.05 1.00 0.13 0.18 -0.12 -0.05 0.07 0.19 -0.12 0.10 0.06 -0.15
chol 0.21 -0.20 -0.07 0.13 1.00 0.01 -0.15 -0.01 0.06 0.05 0.00 0.09 0.10 -0.08
fbs 0.12 0.05 0.10 0.18 0.01 1.00 -0.08 -0.01 0.02 0.00 -0.06 0.14 -0.03 -0.03
restecg -0.11 -0.06 0.04 -0.12 -0.15 -0.08 1.00 0.04 -0.07 -0.06 0.09 -0.08 -0.01 0.13
Simple 2 Python 3 (ipykernel) | Idle
Mode: Edit Ln 5, Col 11 Data Minning