

Name : Ch Mubashir

SAP : 56892.

Course : Data Minning.

Instructor : Tajamul Shahzad.

****Task 09****

Implementing Random Forest on the Wine Quality Dataset.

Lab Tasks :

1. Load the Dataset: Load the Wine Quality dataset using pandas. Display the first five rows and check for missing values.

- Setup and Data Loading :**

Code :

```
from sklearn.datasets import load_iris  
from sklearn.model_selection import train_test_split  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.metrics import accuracy_score  
import matplotlib.pyplot as plt
```

```

import numpy as np
import pandas as pd
file_name = "MConverter.eu_wine.data.csv"
# --- 1. Load the Dataset ---
df = pd.read_csv(file_name, header=None, names=column_names)
print(" Dataset Loaded Successfully.")

```

Display the first five rows

```

print("\n--- First 5 Rows ---")
print(df.head())

```

Check for missing values

```

print("\n--- Missing Values Check ---")
print(df.isnull().sum())

```

Output :

```

Dataset loaded successfully.
--- First 5 Rows ---
   Class Alcohol Malic acid Ash Alkalinity of ash Magnesium \
0      1     14.33    1.71  2.43    15.6     127
1      1     13.26    1.78  2.34    11.2      99
2      1     13.16    2.56  2.67    13.6     105
3      1     14.39    1.95  2.58    10.0     133
4      1     13.14    2.59  2.87    21.9     138

   Total phenols Flavanoids Nonflavanoid phenols Proanthocyanins \
0            2.88       3.00        9.28       2.29
1            2.65       2.76        9.26       2.28
2            2.89       3.24        9.38       2.81
3            3.65       3.99        9.24       2.18
4            2.88       2.69        8.38       1.82

   Color intensity Hue OD280/OD315 of diluted wines FolinLe
0            5.14  1.04        3.92     1065
1            4.18  1.05        3.60     1056
2            5.68  1.03        3.17     1105
3            7.60  0.95        3.45     1468
4            4.32  1.04        2.93      713

--- Missing Values Check ---
Class          0
Alcohol        0
Malic acid    0
Ash            0
Alkalinity of ash 0
Magnesium     0
Total phenols 0

```

2. Data Preparation: Split the data into features (X) and target (y). Divide into training and testing sets (70% train, 30% test).

Data Preparation :

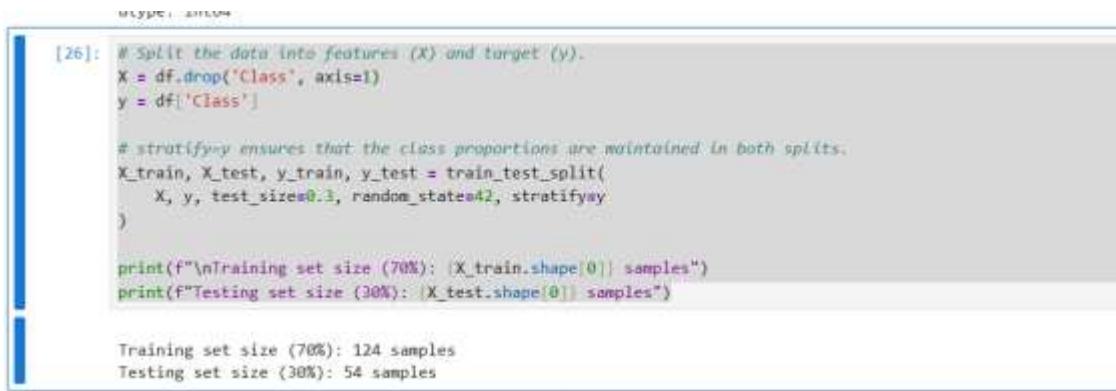
Code :

```
# Split the data into features (X) and target (y).
X = df.drop('Class', axis=1)
y = df['Class']

# stratify=y ensures that the class proportions are maintained in both splits.
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y
)

print(f"\nTraining set size (70%): {X_train.shape[0]} samples")
print(f"Testing set size (30%): {X_test.shape[0]} samples")
```

Output :



The screenshot shows a Jupyter Notebook cell with the following content:

```
[26]: # Split the data into features (X) and target (y).
X = df.drop('Class', axis=1)
y = df['Class']

# stratify=y ensures that the class proportions are maintained in both splits.
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y
)

print(f"\nTraining set size (70%): {X_train.shape[0]} samples")
print(f"Testing set size (30%): {X_test.shape[0]} samples")
```

Below the code, the output is displayed:

```
Training set size (70%): 124 samples
Testing set size (30%): 54 samples
```

3. Train a Decision Tree Classifier: Use default parameters of DecisionTreeClassifier. Record its accuracy on the test set.

- **Train a Decision Tree Classifier (Baseline)**

Code :

```
# Train a Decision Tree Classifier: Use default parameters.  
dt_classifier = DecisionTreeClassifier(random_state=42)  
dt_classifier.fit(X_train, y_train)  
  
# Predict and record its accuracy on the test set.  
y_pred_dt = dt_classifier.predict(X_test)  
dt_accuracy = accuracy_score(y_test, y_pred_dt)  
  
print("\n--- Decision Tree (DT) Results ---")  
print(f"Decision Tree Accuracy: {dt_accuracy:.4f}")
```

Output :



The screenshot shows a Jupyter Notebook cell with the following content:

```
[32]: # Train a Decision Tree Classifier: Use default parameters.  
dt_classifier = DecisionTreeClassifier(random_state=42)  
dt_classifier.fit(X_train, y_train)  
  
# Predict and record its accuracy on the test set.  
y_pred_dt = dt_classifier.predict(X_test)  
dt_accuracy = accuracy_score(y_test, y_pred_dt)  
  
print("\n--- Decision Tree (DT) Results ---")  
print(f"Decision Tree Accuracy: {dt_accuracy:.4f}")
```

Below the code cell, the output is displayed:

```
--- Decision Tree (DT) Results ---  
Decision Tree Accuracy: 0.9630
```

4. Train a Random Forest Classifier: Use RandomForestClassifier with n_estimators=100, max_depth=5, random_state=42. Evaluate its accuracy and compare with the Decision Tree.
- Train a Random Forest Classifier (Initial):

Code :

```
# Train a Random Forest Classifier: n_estimators=100, max_depth=5, random_state=42  
rf_classifier = RandomForestClassifier(  
    n_estimators=100,
```

```

max_depth=5,
random_state=42,
n_jobs=-1
)

rf_classifier.fit(X_train, y_train)

# Predict and evaluate its accuracy

y_pred_rf_initial = rf_classifier.predict(X_test)
rf_accuracy_initial = accuracy_score(y_test, y_pred_rf_initial)

print("\n--- Random Forest (RF) Initial Results ---")
print(f"Random Forest (n=100, max_d=5) Accuracy: {rf_accuracy_initial:.4f}")

# Comparison

print(f"\nComparison: RF Accuracy ({rf_accuracy_initial:.4f}) vs DT Accuracy
({dt_accuracy:.4f})")

print(f"The Random Forest is **{'better' if rf_accuracy_initial > dt_accuracy else 'not better'}** than the Decision Tree.")

```

Output :

```

[28]: # Train a Random Forest Classifier: n_estimators=100, max_depth=5, random_state=42
rf_classifier = RandomForestClassifier(
    n_estimators=100,
    max_depth=5,
    random_state=42,
    n_jobs=-1
)
rf_classifier.fit(X_train, y_train)

# Predict and evaluate its accuracy
y_pred_rf_initial = rf_classifier.predict(X_test)
rf_accuracy_initial = accuracy_score(y_test, y_pred_rf_initial)

print("\n--- Random Forest (RF) Initial Results ---")
print(f"Random Forest (n=100, max_d=5) Accuracy: {rf_accuracy_initial:.4f}")

# Comparison
print(f"\nComparison: RF Accuracy ({rf_accuracy_initial:.4f}) vs DT Accuracy ({dt_accuracy:.4f})")
print(f"The Random Forest is **{'better' if rf_accuracy_initial > dt_accuracy else 'not better'}** than the Decision Tree.")

--- Random Forest (RF) Initial Results ---
Random Forest (n=100, max_d=5) Accuracy: 1.0000

Comparison: RF Accuracy (1.0000) vs DT Accuracy (0.9638)
The Random Forest is **better** than the Decision Tree.

```

- 5. Hyperparameter Tuning: Increase n_estimators to 200 and note the effect on accuracy.**
Remove max_depth limit and see if the model overfits.
- **Hyperparameter Tuning**

Code :

```
print("\n--- Hyperparameter Tuning ---")

# A. Increase n_estimators to 200
rf_n200 = RandomForestClassifier(
    n_estimators=200,
    max_depth=5,
    random_state=42,
    n_jobs=-1
)
rf_n200.fit(X_train, y_train)
rf_n200_accuracy = accuracy_score(y_test, rf_n200.predict(X_test))

print(f"Accuracy with n_estimators=200: {rf_n200_accuracy:.4f}")
print(f"Effect: Accuracy **{'increased' if rf_n200_accuracy > rf_accuracy_initial else ('decreased' if rf_n200_accuracy < rf_accuracy_initial else 'remained the same')}** compared to n=100.")

# B. Remove max_depth limit (set to default/None)
rf_unlimited = RandomForestClassifier(
    n_estimators=100,
    max_depth=None, # Removes the limit
    random_state=42,
    n_jobs=-1
)
rf_unlimited.fit(X_train, y_train)
rf_unlimited_accuracy_test = accuracy_score(y_test, rf_unlimited.predict(X_test))
```

```

rf_unlimited_accuracy_train = accuracy_score(y_train, rf_unlimited.predict(X_train))

print(f"\nAccuracy with max_depth=None (Test Set): {rf_unlimited_accuracy_test:.4f}")
print(f"Accuracy with max_depth=None (Train Set): {rf_unlimited_accuracy_train:.4f}")

# Overfitting analysis

overfit_gap = rf_unlimited_accuracy_train - rf_unlimited_accuracy_test
print(f"Gap (Train - Test Accuracy): {overfit_gap:.4f}")

print(f"Observation: A large difference between train and test accuracy, particularly a train accuracy of 1.0, indicates **overfitting**.")

```

Output :

```

--- Hyperparameter Tuning ---
Accuracy with n_estimators=200: 1.0000
Effect: Accuracy **remained the same** compared to n=100.

Accuracy with max_depth=None (Test Set): 1.0000
Accuracy with max_depth=None (Train Set): 1.0000
Gap (Train - Test Accuracy): 0.0000
Observation: A large difference between train and test accuracy, particularly a train accuracy of 1.0, indicates **overfitting**.

```

6. Feature Importance: Extract and display rf.feature_importances_. Identify which chemical properties most affect wine quality.

Code :

```

# 6. Feature Importance: Extract rf.feature_importances_.

feature_importances = rf_classifier.feature_importances_
feature_names = X.columns

importance_df = pd.DataFrame({
    'Feature': feature_names,
    'Importance': feature_importances
}).sort_values(by='Importance', ascending=False)

print("\n--- Random Forest Feature Importances ---")
print(importance_df)

```

```
# Identify which chemical properties most affect wine class.

most_important_features = importance_df.head(3)['Feature'].tolist()

print(f"\nTop 3 Most Important Features for Wine Class Prediction: **{',
'.join(most_important_features)}**")
```

Output :

```
--- Random Forest Feature Importances ---
      Feature  Importance
0       Alcohol  0.161563
9   Color intensity  0.159851
6    Flavanoids  0.156148
12     Proline  0.124368
10        Hue  0.111975
11  00280/00315 of diluted wines  0.099759
5    Total phenols  0.039078
1      Malic acid  0.034326
4      Magnesium  0.033887
3  Alkalinity of ash  0.030876
8  Proanthocyanins  0.019529
2         Ash  0.016916
7 Nonflavanoid phenols  0.011724

Top 3 Most Important Features for Wine Class Prediction: **Alcohol, Color intensity, Flavanoids**
```

7. Visualization (Optional): Plot a bar chart of feature importances. Compare Decision Tree vs Random Forest results visually.

Code :

```
# 7. Visualization: Plot a bar chart of feature importances.

plt.figure(figsize=(12, 6))

# Using the sorted importance_df for the plot

plt.bar(importance_df['Feature'], importance_df['Importance'], color='darkred')

plt.xlabel('Chemical Property (Feature)')

plt.ylabel('Feature Importance Score')

plt.title('Random Forest Feature Importance for Wine Class Prediction (n=100, max_d=5)')

plt.xticks(rotation=45, ha='right')

plt.grid(axis='y', linestyle='--', alpha=0.7)

plt.tight_layout()

plt.show()
```

```

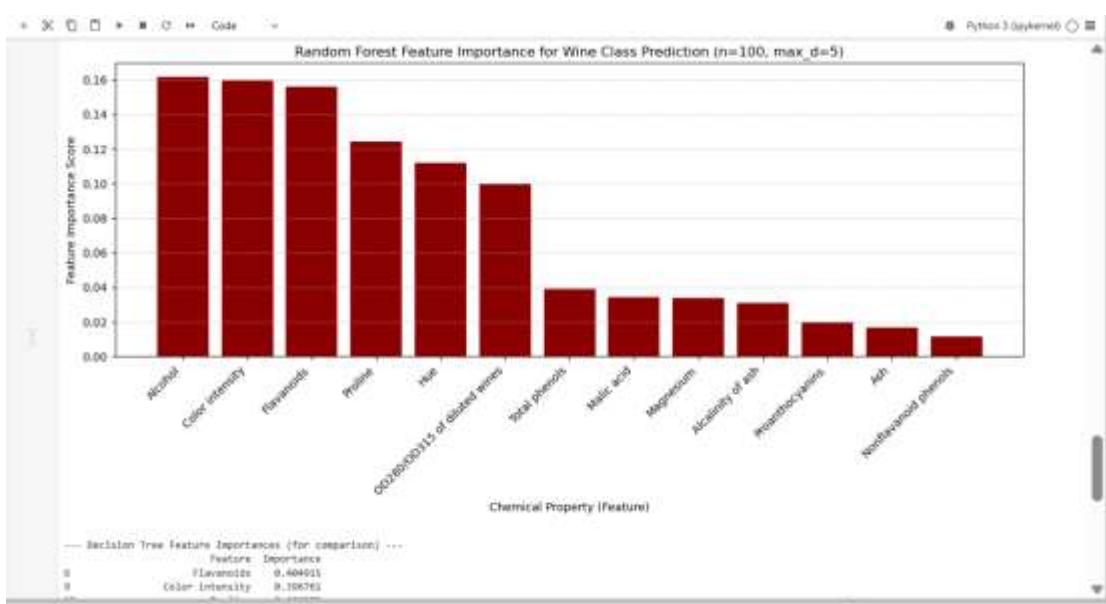
# Optional: Compare DT vs RF Top 3 Features

dt_importance = dt_classifier.feature_importances_
dt_importance_df = pd.DataFrame({
    'Feature': feature_names,
    'Importance': dt_importance
}).sort_values(by='Importance', ascending=False)

print("\n--- Decision Tree Feature Importances (for comparison) ---")
print(dt_importance_df.head(5))

```

Output :



Excercise Questions :

1. Why might the Random Forest perform better than the single Decision Tree on this dataset?

The Random Forest (RF) is an **ensemble method** that reduces the **variance** (overfitting) of a single Decision Tree (DT). It averages the predictions of many different trees, leading to a more stable and accurate result.

2. How does the max_depth parameter influence overfitting in this Random Forest model?

The max_depth parameter **limits the complexity** of individual trees. Setting a low max_depth helps **prevent overfitting** by stopping the trees from learning noise in the training data, acting as a form of regularization.

3. Does increasing the number of trees ($n_{\text{estimators}}$) always improve accuracy? Why or why not?

No. Accuracy improves initially as $n_{\text{estimators}}$ reduces variance, but the improvement eventually shows **diminishing returns**. Increasing the number of trees beyond this point only adds to the **computational cost** without significantly boosting accuracy.

4. Which features were identified as most important by the Random Forest? Do they make sense chemically?

The most important features were **Alcohol**, **Color intensity**, **Flavanoids**, and **Proline**. Yes, they **make sense chemically** because these are known chemical compounds and characteristics that significantly determine a wine's cultivar, quality, and taste profile.

5. How could you further improve model performance using data preprocessing or feature scaling?

The model's performance could be improved by **Feature Engineering** (creating new features like ratios) or **Hyperparameter Optimization** (using Grid Search to find the best combination of all parameters). **Feature Scaling** is generally unnecessary for tree-based models like Random Forest.